

Licensed under the Apache License, Version 2.0 (the "License");



[Run in Google Colab](#)



[View source on GitHub](#)

▼ TensorFlow Hub

[TensorFlow Hub](#) is an online repository of already trained TensorFlow models that you can use. These models can be used for Transfer Learning.

Transfer learning is a process where you take an existing trained model, and extend it to do additional work. This is done by keeping the model unchanged, while adding and retraining the final layers, in order to get a different set of possible outputs.

Here, you can see all the models available in [TensorFlow Module Hub](#).

Before starting this Colab, you should reset the Colab environment by selecting Runtime -> Reset all runtimes...

▼ Imports

This Colab will require us to use some things which are not yet in official releases of TensorFlow. So below, we're installing a nightly version of TensorFlow as well as TensorFlow Hub.

This will switch your installation of TensorFlow in Colab to this TensorFlow version. Once you are finished with the tutorial, you can switch back to the latest stable release of TensorFlow by doing selecting Runtime -> Reset all runtimes... in the Colab interface to reset the Colab environment to its original state.

```
!pip install tf-nightly-gpu
!pip install "tensorflow_hub==0.4.0"
!pip install -U tensorflow_datasets
```



```

Requirement already satisfied: tf-nightly-gpu in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: tf-estimator-nightly in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: tb-nightly<1.15.0a0,>=1.14.0a0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: keras-preprocessing>=1.0.5 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: six>=1.10.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: google-pasta>=0.1.6 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: absl-py>=0.7.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: numpy<2.0,>=1.14.5 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: grpcio>=1.8.6 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: wheel>=0.26 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: keras-applications>=1.0.8 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: protobuf>=3.6.1 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: gast>=0.2.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: wrapt>=1.11.1 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: astor>=0.6.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: h5py in /usr/local/lib/python3.6/dist-packages (from keras)
Requirement already satisfied: tensorflow_hub==0.4.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: six>=1.10.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: protobuf>=3.4.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: numpy>=1.12.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already up-to-date: tensorflow_datasets in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: psutil in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: protobuf>=3.6.1 in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: requests in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: dill in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: six in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: tqdm in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: wrapt in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: absl-py in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: promise in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: numpy in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: future in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: termcolor in /usr/local/lib/python3.6/dist-packages (from tensorflow)
Requirement already satisfied, skipping upgrade: tensorflow-metadata in /usr/local/lib/python3.6/dist-packages (from tensorflow)

```

Some normal imports we've seen before. The new one is importing tensorflow_hub which was installed above, and we'll make heavy use of.

```

from __future__ import absolute_import, division, print_function, unicode_literals

import numpy as np
import matplotlib.pyplot as plt

import tensorflow as tf
tf.enable_eager_execution()

import tensorflow_hub as hub
import tensorflow_datasets as tfds

from tensorflow.keras import layers

```

```

import logging
logger = tf.get_logger()
logger.setLevel(logging.ERROR)

```

▼ TODO: Download the Flowers Dataset using TensorFlow

In the cell below you will download the Flowers dataset using TensorFlow Datasets. If you look at the [TensorFlow Datasets](#) you will see that the name of the Flowers dataset is `tf_flowers`. You can also see that this dataset is only split into training and validation sets, so you will therefore have to use `tfds.splits` to split this training set into to a `training_set` and a `validation_set`. The `tfds.splits` function that 70 corresponds to the `training_set` and 30 to the `validation_set`. Then load the `tf_flowers` dataset using the `tfds.load` function. The `tfds.load` function uses all the parameters you need, and also make sure it returns the dataset info, so you can print information about the datasets.

```
splits = tfds.Split.TRAIN.subsplit([70, 30])

(training_set, validation_set), dataset_info = tfds.load('tf_flowers', with_info=True, as_supervised=True)
```

▼ TODO: Print Information about the Flowers Dataset

Now that you have downloaded the dataset, use the dataset info to print the number of classes in the dataset, and also print the number of images in the training and validation sets.

```
num_classes = dataset_info.features['label'].num_classes

num_training_examples = 0
num_validation_examples = 0

for example in training_set:
    num_training_examples += 1

for example in validation_set:
    num_validation_examples += 1

print('Total Number of Classes: {}'.format(num_classes))
print('Total Number of Training Images: {}'.format(num_training_examples))
print('Total Number of Validation Images: {} \n'.format(num_validation_examples))
```

```
☞ Total Number of Classes: 5
   Total Number of Training Images: 2590
   Total Number of Validation Images: 1080
```

The images in the Flowers dataset are not all the same size.

```
for i, example in enumerate(training_set.take(5)):
    print('Image {} shape: {} label: {}'.format(i+1, example[0].shape, example[1]))
```

```
☞ Image 1 shape: (335, 500, 3) label: 2
   Image 2 shape: (257, 320, 3) label: 0
   Image 3 shape: (213, 320, 3) label: 0
   Image 4 shape: (213, 320, 3) label: 2
   Image 5 shape: (240, 320, 3) label: 0
```

▼ TODO: Reformat Images and Create Batches

In the cell below create a function that reformats all images to the resolution expected by MobileNet v2 (224, 224). The function should take in an image and a label as arguments and should return the new image and corresponding label. Finally, create training and validation batches of size 32.

```
IMAGE_RES = 224

def format_image(image, label):
    image = tf.image.resize(image, (IMAGE_RES, IMAGE_RES))/255.0
    return image, label

BATCH_SIZE = 32

train_batches = training_set.shuffle(num_training_examples//4).map(format_image).batch(BATCH_SIZE)
validation_batches = validation_set.map(format_image).batch(BATCH_SIZE).prefetch(1)
```

▼ Do Simple Transfer Learning with TensorFlow Hub

Let's now use TensorFlow Hub to do Transfer Learning. Remember, in transfer learning we reuse parts of an already trained model (the feature extractor), change the final layer, or several layers, of the model, and then retrain those layers on our own dataset.

TODO: Create a Feature Extractor

In the cell below create a feature_extractor using MobileNet v2. Remember that the partial model from TensorFlow Hub (the feature extractor, i.e., the model without the final classification layer) is called a feature vector. Go to the [TensorFlow Hub documentation](https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/2) to see a list of available feature vectors. Read the documentation and get the corresponding URL for the feature vector. Finally, create a feature_extractor by using hub.KerasLayer with the correct input_shape parameter.

```
URL = "https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/2"
feature_extractor = hub.KerasLayer(URL,
                                   input_shape=(IMAGE_RES, IMAGE_RES, 3))
```

▼ TODO: Freeze the Pre-Trained Model

In the cell below freeze the variables in the feature extractor layer, so that the training only modifies the final classification layer.

```
feature_extractor.trainable = False
```

▼ TODO: Attach a classification head

In the cell below create a tf.keras.Sequential model, and add the pre-trained model and the new classification layer. The new classification layer must have the same number of classes as our Flowers dataset. Finally print a summary of the model.

```
model = tf.keras.Sequential([
    feature_extractor,
    layers.Dense(num_classes, activation='softmax')
])

model.summary()
```

🔗 Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
keras_layer (KerasLayer)	(None, 1280)	2257984
=====		
dense (Dense)	(None, 5)	6405
=====		
Total params: 2,264,389		
Trainable params: 6,405		
Non-trainable params: 2,257,984		
=====		

▼ TODO: Train the model

In the cell bellow train this model like any other, by first calling `compile` and then followed by `fit`. Make sure you when applying both methods. Train the model for only 6 epochs.

```
model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])

EPOCHS = 6

history = model.fit(train_batches,
                    epochs=EPOCHS,
                    validation_data=validation_batches)
```

🔗 Epoch 1/6
81/81 [=====] - 20s 248ms/step - loss: 0.9136 - acc: 0.6811
Epoch 2/6
81/81 [=====] - 9s 113ms/step - loss: 0.5218 - acc: 0.8483 -
Epoch 3/6
81/81 [=====] - 9s 114ms/step - loss: 0.4321 - acc: 0.8857 -
Epoch 4/6
81/81 [=====] - 9s 114ms/step - loss: 0.3792 - acc: 0.9054 -
Epoch 5/6
81/81 [=====] - 9s 114ms/step - loss: 0.3414 - acc: 0.9208 -
Epoch 6/6
81/81 [=====] - 9s 114ms/step - loss: 0.3119 - acc: 0.9274 -

You can see we get ~88% validation accuracy with only 6 epochs of training, which is absolutely awesome. This the model we created in the previous lesson, where we were able to get ~76% accuracy with 80 epochs of training. The difference is that MobileNet v2 was carefully designed over a long time by experts, then trained on a massive dataset.

▼ TODO: Plot Training and Validation Graphs.

In the cell below, plot the training and validation accuracy/loss graphs.

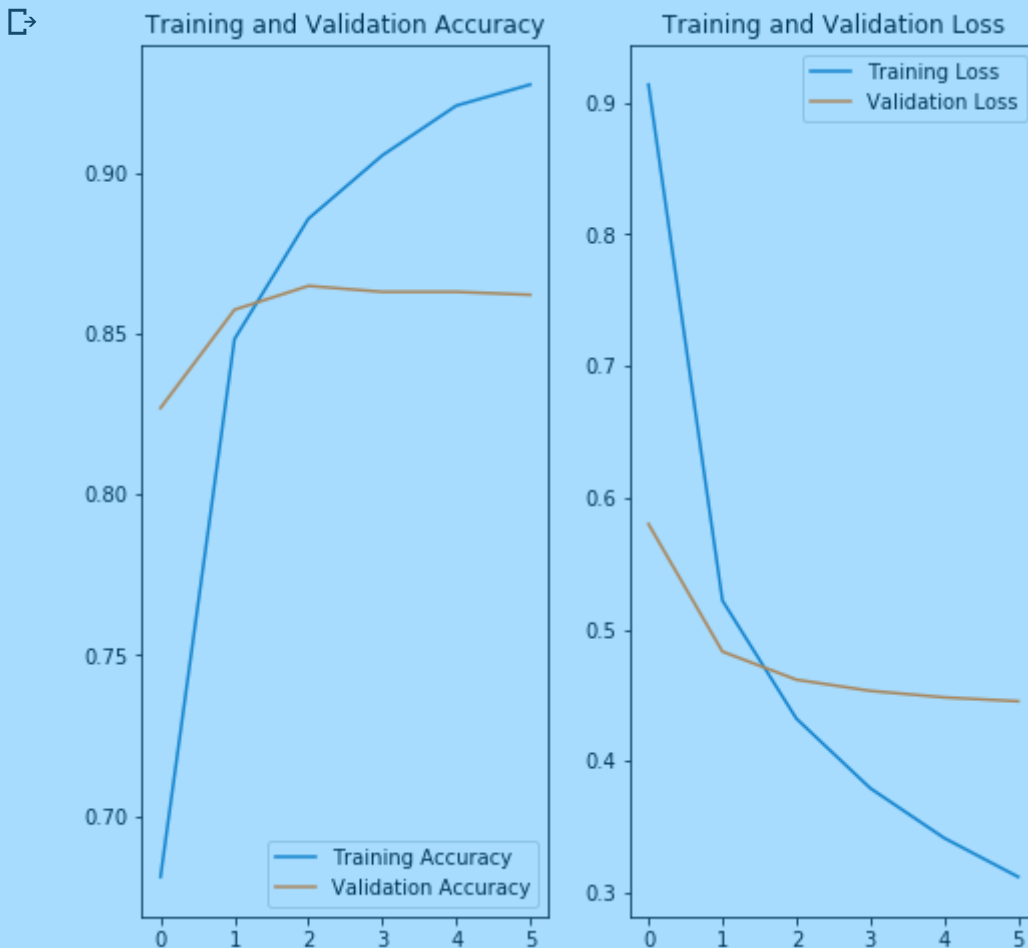
```
acc = history.history['acc']
val_acc = history.history['val_acc']

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(EPOCHS)
```

```
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



What is a bit curious here is that validation performance is better than training performance, right from the start. One reason for this is that validation performance is measured at the end of the epoch, but training performance is measured across the epoch.

The bigger reason though is that we're reusing a large part of MobileNet which is already trained on Flower images.

▼ TODO: Check Predictions

In the cell below get the label names from the dataset info and convert them into a NumPy array. Print the array to see the correct label names.

```
class_names = np.array(dataset_info.features['label'].names)
```

```
print(class_names)
```

```
['dandelion' 'daisy' 'tulips' 'sunflowers' 'roses']
```

▼ TODO: Create an Image Batch and Make Predictions

In the cell below, use the `next()` function to create an `image_batch` and its corresponding `label_batch`. Convert `image_batch` and `label_batch` to numpy arrays using the `.numpy()` method. Then use the `.predict()` method to run the image model and make predictions. Then use the `np.argmax()` function to get the indices of the best prediction for each image. Finally, use the `class_names` list to get the names of the classes corresponding to the indices of the best predictions to class names.

```
image_batch, label_batch = next(iter(train_batches))
```

```
image_batch = image_batch.numpy()
label_batch = label_batch.numpy()
```

```
predicted_batch = model.predict(image_batch)
predicted_batch = tf.squeeze(predicted_batch).numpy()
```

```
predicted_ids = np.argmax(predicted_batch, axis=-1)
predicted_class_names = class_names[predicted_ids]
```

```
print(predicted_class_names)
```

```
['dandelion' 'tulips' 'sunflowers' 'tulips' 'daisy' 'sunflowers' 'daisy'
 'tulips' 'dandelion' 'tulips' 'roses' 'tulips' 'roses' 'tulips' 'roses'
 'sunflowers' 'roses' 'dandelion' 'tulips' 'roses' 'daisy' 'sunflowers'
 'dandelion' 'daisy' 'sunflowers' 'tulips' 'daisy' 'sunflowers'
 'sunflowers' 'tulips' 'dandelion' 'tulips']
```

▼ TODO: Print True Labels and Predicted Indices

In the cell below, print the true labels and the indices of predicted labels.

```
print("Labels: ", label_batch)
print("Predicted labels: ", predicted_ids)
```

```
Labels:  [0 3 3 3 1 3 1 2 0 2 4 2 4 2 4 3 4 0 2 4 1 3 0 1 3 2 1 3 3 4 0 2]
Predicted labels:  [0 2 3 2 1 3 1 2 0 2 4 2 4 2 4 3 4 0 2 4 1 3 0 1 3 2 1 3 3 2 0 2]
```

▼ Plot Model Predictions

```
plt.figure(figsize=(10,9))
for n in range(30):
    plt.subplot(6,5,n+1)
    plt.imshow(image_batch[n])
    color = "blue" if predicted_ids[n] == label_batch[n] else "red"
    plt.title(predicted_class_names[n].title(), color=color)
    plt.axis('off')
_ = plt.suptitle("Model predictions (blue: correct, red: incorrect)")
```

Model predictions (blue: correct, red: incorrect)

Dandelion



Sunflowers



Roses



Sunflowers




Daisy



Tulips



Tulips



Daisy



Tulips



Roses



Sunflowers



Daisy



Sunflowers



Tulips



Roses



Dandelion



Dandelion



Sunflowers



Tulips



Dandelion



Tulips



Tulips



Daisy




Sunflowers



Daisy



Tulips



Roses



Roses



Sunflowers



Tulips

